

CS 615 – Deep Learning

Artificial Neurons

Slides adapted from material created by E. Alpaydin
Prof. Mordohai, Prof. Greenstadt, Pattern Classification (2nd Ed.),
Pattern Recognition and Machine Learning

Objectives

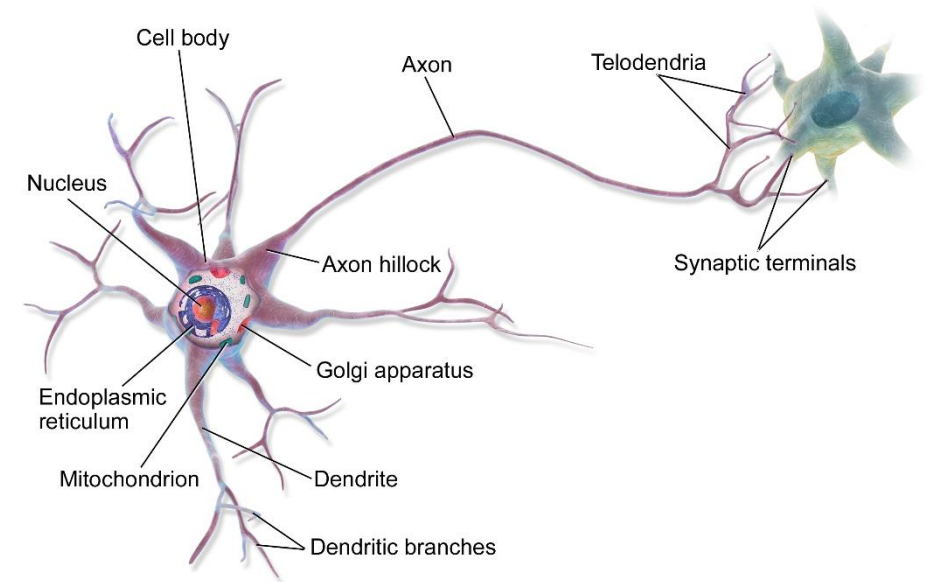
- Artificial Neurons
- Evaluating a classifier
- Logistic activation function
- Log likelihood objective function
- Softmax activation function with cross entropy objective function

Artificial Neurons

- In the previous material (on **linear regression**) we discussed the idea of estimating an output value as $\hat{y} = g(x) = x\theta$ where θ is a vector of parameters.
- And to learn the values of these parameters, we can start with an initial guess, and iteratively update them based on the gradient of the least squared error.
- But what if the output was a binary value?
 - Or a probability?
- Then we could use it for binary classification!

Artificial Neurons

- A simple “network” that takes the weighed sum of the inputs, $x\theta$ and applies some function to it is called an *artificial neuron*
- These are the building blocks of artificial and deep networks.



Artificial Neurons

- Technically a linear regression module is an artificial neuron.
- The function that processes $x\theta$ is referred to as the **activation function**
- For linear regression this was the linear activation function:

$$g(x) = x\theta$$

- At this time we're also going to move to the notation $g(x|\theta)$ for reasons seen later.
- A **non-linear activation function** is the threshold/step function:
$$g(x) = x\theta > t$$
- An artificial neuron that uses this is referred to as a **perceptron**.

Artificial Neurons for Classification

- Now we want to use artificial neurons for classification.
- In doing so we'll have several decisions:
 1. What is our **activation function**?
 2. What is our **objective function**?
- Both of these give rise to a **particular gradient rule** used to learn the model's parameters.
- Let's look at a few common combinations....

Finding the Weights

- Let's start with the squared error objective function.
- Recall that given some function $g(x|\theta)$, **for a single observation**, (x, y) , we can generalize the squared error as:

$$J = (y - \hat{y})^2 = (y - g(x|\theta))^2$$

- The gradient of this, with respect to a single parameter θ_i , is then:

$$\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$$

Gradient of SE

$$\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$$

- So we obviously need to find the partial derivative of g with respect to θ_i
 - Note: With linear regression, $g(x|\theta) = x\theta$ and therefore $\frac{\partial}{\partial \theta_i} g(x|\theta) = x_i$
 - Which results in a gradient of $\frac{\partial J}{\partial \theta_i} = -2x_i(y - x\theta)$
- So what is $\frac{\partial}{\partial \theta_i} g(x|\theta)$ when $g(x|\theta) = x\theta > 0$?

Update Rule

$$\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$$

- So what is $\frac{\partial}{\partial \theta_i} g(x|\theta)$ when $g(x|\theta) = x\theta > 0$?
- There's basically two cases:
 - Case 1: When $x\theta > 0$
 - $\frac{\partial}{\partial \theta_i} g(x|\theta) = x_i$
 - Case 2: When $x\theta \leq 0$
 - $\frac{\partial}{\partial \theta_i} g(x|\theta) = 0$
- Therefore:

$$\frac{\partial}{\partial \theta_i} (x\theta > 0) = \begin{cases} x_i & \text{if } x\theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Update Rule

$$\frac{\partial}{\partial \theta_i} (x\theta > 0) = \begin{cases} x_i & \text{if } x\theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Plugging this back into $\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$ we get our gradient update rule for SE with perceptrons as:

$$\frac{\partial J}{\partial \theta_i} = \begin{cases} -2x_i(y - 1) & \text{if } x\theta > 0 \\ \cancel{-2x_i y} & \text{otherwise} \end{cases} \quad \text{0}$$

- Which we can also write as:

$$\frac{\partial J}{\partial \theta_i} = \begin{cases} 2x_i(1 - y) & \text{if } x\theta > 0 \\ \cancel{-2x_i y} & \text{otherwise} \end{cases} \quad \text{0}$$

- We can vectorize the parameters:

$$\frac{\partial J}{\partial \theta} = \begin{cases} \cancel{2x^T(1 - y)} & \text{if } x\theta > 0 \\ \cancel{-2x^T y} & \text{otherwise} \end{cases}$$

Logistic Activation Function

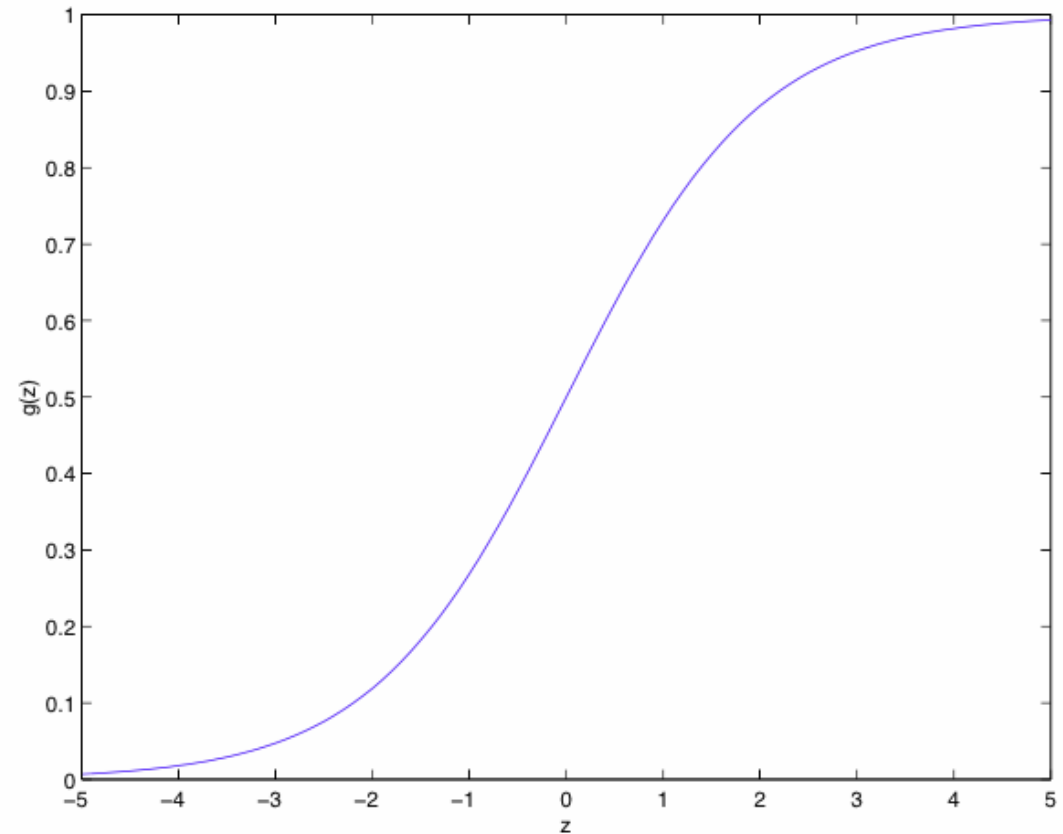
- Let's take a look at another activation function.
- An activation function this is non-linear, but differentiable, is the *logistic function* (sometimes known as the **sigmoid** or logit function):

$$g(x|\theta) = \frac{1}{1 + e^{-x\theta}}$$

Logistic Function

- The logistic function, defined as $g(z) = \frac{1}{1+e^{-z}}$ tends towards 0 as z decreases and tends towards 1 as z increases
- For our purposes $z = x\theta$ and therefore our **activation function** is:

$$g(x|\theta) = \frac{1}{1 + e^{-x\theta}}$$



Logistic Function

- Let's see how we can use the logistic function, with the least squared error objective function, to define our gradient rule!
- Recall for the LSE:

$$J = (y - \hat{y})^2 = (y - g(x|\theta))^2$$

- And the gradient of this is:

$$\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$$

- What is $\frac{\partial}{\partial \theta_i} g(x|\theta)$ for the logistic activation function?

Derivative of Logistic Function

- $\frac{\partial}{\partial \theta_i} g(x|\theta) = \frac{\partial}{\partial \theta_i} \left(\frac{1}{1+e^{-x\theta}} \right) = \frac{\partial}{\partial \theta_i} \left(1 + e^{-x\theta} \right)^{-1}$
- $= -1(0 - x_i e^{-x\theta}) (1 + e^{-x\theta})^{-2} = \frac{x_i e^{-x\theta}}{(1+e^{-x\theta})^2}$
- $= x_i \frac{1}{1+e^{-x\theta}} \frac{e^{-x\theta}}{1+e^{-x\theta}}$
- $= x_i g(x|\theta) (1 - g(x|\theta))$

Gradient Rule w/ Logistic Function

- So when $g(x|\theta)$ is the logistic function, its partial derivative is:

$$\frac{\partial}{\partial \theta_i} g(x|\theta) = x_i g(x|\theta) (1 - g(x|\theta))$$

- Plugging this back into $\frac{\partial J}{\partial \theta_i} = -2 \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) (y - g(x|\theta))$ we get our gradient update rule for the SE of a perceptrons with a logistic activation function to be:

$$\begin{aligned} \frac{\partial J}{\partial \theta_i} &= -2 \left(x_i g(x|\theta) (1 - g(x|\theta)) \right) (y - g(x|\theta)) \\ &= -2 x_i \hat{y} (1 - \hat{y}) (y - \hat{y}) \end{aligned}$$

Gradient Rule w/ Logistic Function

$$\frac{\partial J}{\partial \theta_i} = -2x_i \hat{y}(1 - \hat{y})(y - \hat{y})$$

- Now let's try to vectorize!
- Vectorizing over all parameters is now:

$$\frac{\partial J}{\partial \theta} = -2x^T \hat{y}(1 - \hat{y})(y - \hat{y})$$

- And our batch gradient update is:

$$\frac{\partial J}{\partial \theta} = -\frac{2}{N} X^T \hat{Y}(1 - \hat{Y})^T (Y - \hat{Y})$$

Gradient Rule w/ Logistic Function

$$\frac{\partial J}{\partial \theta} = -\frac{2}{N} X^T \hat{Y} (1 - \hat{Y})^T (Y - \hat{Y})$$

- Ok, at least this isn't piecewise/conditional!
- But it's not too pretty.
- There are other objective functions that are (often) better suited for classification.
- One such function is called the *log likelihood*

Log Likelihood

- Given a observation (x, y) , we can compute the **likelihood** that we are correct as

$$\ell = (\hat{y})^y (1 - \hat{y})^{(1-y)} = (g(x|\theta))^y (1 - g(x|\theta))^{(1-y)}$$

- Our technique will of course be to find the gradient of this with respect to one of the parameters.
- However, this can be made easier if we first take the *log* of this.
- Recall
 - The log of a product is the sum of its logs: $\ln(mn) = \ln(m) + \ln(n)$
 - The log of an exponent is the exponent times the log of its base: $\ln(a^x) = x\ln(a)$

- Therefore our log likelihood **objective function** is:

$$J = \ln(\ell) = y \ln(g(x|\theta)) + (1 - y) \ln(1 - g(x|\theta))$$

To Maximum Likelihood

$$J = y \ln(g(x|\theta)) + (1 - y) \ln(1 - g(x|\theta))$$

- First off, a reminder...

$$\frac{\partial}{\partial x} (\ln x) = \frac{1}{x} \cdot \frac{\partial}{\partial x} (x)$$

- Therefore

$$\frac{\partial J}{\partial \theta_i} = \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) \frac{y}{g(x|\theta)} + \left(\frac{\partial}{\partial \theta_i} (1 - g(x|\theta)) \right) \frac{(1 - y)}{1 - g(x|\theta)}$$

- And of course $\frac{\partial}{\partial \theta_i} g(x|\theta)$ depends on what our activation function $g(x|\theta)$ is.

To Maximum Likelihood

$$\frac{\partial J}{\partial \theta_i} = \left(\frac{\partial}{\partial \theta_i} g(x|\theta) \right) \frac{y}{g(x|\theta)} + \left(\frac{\partial}{\partial \theta_i} (1 - g(x|\theta)) \right) \frac{(1 - y)}{1 - g(x|\theta)}$$

- If we choose the logistic function to be our activation function, its partial derivative is:

$$\frac{\partial}{\partial \theta_i} g(x|\theta) = x_i g(x|\theta) (1 - g(x|\theta))$$

- And therefore our gradient is just:

$$\frac{\partial J}{\partial \theta_i} = x_i (y - g(x|\theta)) = x_i (y - \hat{y})$$

To Maximum Likelihood

$$\frac{\partial J}{\partial \theta_i} = x_i(y - \hat{y})$$

- Vectorizing this for all parameters we have

$$\frac{\partial J}{\partial \theta} = x^T(y - \hat{y})$$

- And our batch gradient is just:

$$\frac{\partial J}{\partial \theta} = \frac{1}{N} X^T(Y - \hat{Y})$$

- What we just derived (maximizing the log likelihood based on the output of a logistic function) is called *logistic regression*
 - Which is confusing since it's actually used for *classification*, not regression.
 - Nevertheless...

Gradient Ascent Rule

- One last thing!
- Since we want to *maximize* the log likelihood, we want to ***add*** some amount of the gradient to the parameters.

$$\theta := \theta + \eta \frac{\partial J}{\partial \theta}$$

Evaluation

- How do we evaluate when we're doing classification?
- I suppose we could use the RMSE, but that may make less sense.
- We could just count how often we predict the correct class
 - We call this *accuracy*.
 - Let Y_i be the true class, and \hat{Y}_i be the predicted class for observation i .

$$accuracy = \frac{1}{N} \sum_{i=1}^N (Y_i == \hat{Y}_i)$$

Class Priors

- It's also important to have some sort of baseline accuracy.
- The *class prior* is the likelihood (probability) of a class.
- So for our baseline, we can just use the highest of the class priors.
 - Since naively assigning all test samples that class, would result in that accuracy.

Binary Classification Error Types

- If we're doing binary classification (just two classes), then there's some additional evaluations we can do.
- In binary classification, often we're focused on attempting to "find" one particular class.
 - We refer to this as the positive class.
 - The other data is called the negative class.
- From this we can describe four different possibilities:
 - True positive = Hit
 - True negative = Correct rejection
 - False positive = False Alarm (Type 1 error)
 - False negative = Miss (Type 2 error)

	Predicted positive	Predicted negative	
Positive examples	True positives	False negatives	
Negative examples	False positives	True negatives	

Evaluating your Classifier

- From the four error types, we can establish some binary-classification-specific measurements:
- *Precision* – percentage of things that were classified as positive and actually were positive
 - $Precision = \frac{TP}{TP+FP}$
- *Recall* – the percentage of true positives (*sensitivity*) correctly identified
 - $Recall = \frac{TP}{TP+FN}$
- *f-measure* – The weighted harmonic mean of precision and recall
 - $F_1 = \frac{2*precision*recall}{precision+recall}$

Evaluating your Classifier

- Related to this, we sometimes will look at the true positive rate vs the false positive rate

- The true positive rate is

$$TPR = Recall = \frac{TP}{TP + FN}$$

- The false positive rate is

$$FPR = \frac{FP}{FP + TN}$$

Logistic Regression Example

- Let's classifying whether a person will buy a product or not

Obs. No.	Y	X-Variables							Prev Child Mag	Prev Parent Mag
	Buy	Income	Is Female	Is Married	Has College	Is Professional	(Omitted Variables)			
1	0	24000	1	0	1	1	...	0	0	
2	1	75000	1	1	1	1	...	1	0	
3	0	46000	1	1	0	0	...	0	0	
4	1	70000	0	1	0	1	...	1	0	
5	0	43000	1	0	0	0	...	0	1	
6	0	24000	1	1	0	0	...	0	0	
7	0	26000	1	1	1	0	...	0	0	
8	0	38000	1	1	0	0	...	0	0	
9	0	39000	1	0	1	1	...	0	0	
10	0	49000	0	1	0	0	...	0	0	
.	
.	
.	
654	0	10000	1	0	0	0	...	0	0	
655	1	75000	0	1	0	1	...	0	0	
656	0	72000	0	0	1	0	...	0	0	
657	0	33000	0	0	0	0	...	0	0	
658	0	58000	0	1	1	1	...	0	0	
659	1	49000	1	1	0	0	...	0	0	
660	0	27000	1	1	0	0	...	0	0	
661	0	4000	1	0	0	0	...	0	0	
662	0	40000	1	0	1	1	...	0	0	
663	0	75000	1	1	1	0	...	0	0	
664	0	27000	1	0	0	0	...	0	0	
665	0	22000	0	0	0	1	...	0	0	
666	0	8000	1	1	0	0	...	0	0	
667	1	75000	1	1	1	0	...	0	0	
668	0	21000	0	1	0	0	...	0	0	
669	0	27000	1	0	0	0	...	0	0	
670	0	3000	1	0	0	0	...	0	0	
671	1	75000	1	1	0	1	...	0	0	
672	1	51000	1	1	0	1	...	0	0	
673	0	11000	0	1	0	0	...	0	0	

KidCreative.csv

Logistic Regression Example

- Make some design decisions:
 - Randomize data
 - Use 2/3 training, 1/3 testing
 - Standardize features
 - Add bias feature
 - Initialize parameters to random values in the range $[-1, 1]$
 - Since our equation is based on log likelihood, let's terminate when change in sum of log likelihoods doesn't change more than ϵ
 - Recall the log likelihood of an example being correct is
$$y \ln \left(\frac{1}{1 + e^{-x\theta}} \right) + (1 - y) \ln \left(1 - \frac{1}{1 + e^{-x\theta}} \right)$$
 - But be careful... $\log(0) = -\text{Inf}$. So you might need to deal with this somehow
 - Let's dynamically allow the learning parameter η to adapt!
 - Let's do batch regression

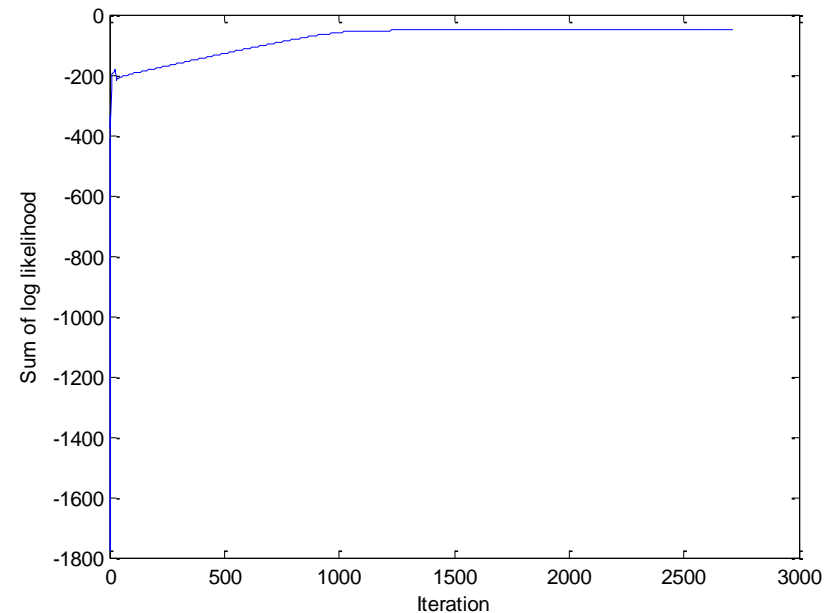
Example

$\theta =$


-12.7663
5.1987
0.9057
0.6136
0.0880
-0.3606
-0.4920
-29.3976
0.2140
-0.0329
0.7355
0.3957
0.0054
0.9061
1.1647
0.1887
0.1881

Choosing Class 1 if $g_{\theta}(x) \geq 0.5$ we get:

Precision: 0.7708
Recall: 0.7551
F-Measure: 0.7629



Example



Multiple Classes

- What if we had multiple classes to decide among?
- There are ways that we can make this decision **using a bunch of binary classifiers**, but let's look at how we can do this directly.

Multi-Class MLE + Logistic Activation

- In order to do multi-class classification, there will be three main changes:
 1. Our target for each observation will now be a *binary vector*. Therefore $Y \in \mathcal{R}^{N \times K}$
 2. Likewise, our estimations will be a vector: $\hat{y} \in \mathcal{R}^{1 \times K}$
 3. And finally, our parameters will now be a $D \times K$ **matrix**
- Given all of this, let's adapt our **maximum** likelihood with logistic activation equations and gradient rules for multi-class classification:

$$\ell_{\theta}(x, y) = \prod_{k=1}^K g(x|\theta_{:,k})^{y_k} (1 - g(x|\theta_{:,k}))^{(1-y_k)}$$

- Taking the log of this we get our objective function:

$$J = \sum_{k=1}^K y_k \ln(g(x|\theta_{:,k})) + (1 - y_k) \ln(1 - g(x|\theta_{:,k}))$$

Multi-Class MLE + Logistic Activation

$$J = \sum_{k=1}^K y_k \ln(g(x|\theta_{:,k})) + (1 - y_k) \ln(1 - g(x|\theta_{:,k}))$$

- Now we need to find the gradient with respect to each $\theta_{i,k}$, i.e. $\frac{\partial J}{\partial \theta_{i,k}}$
- Fortunately, the only term of the summation that will contribute is when $k = k$.
- Therefore we have:

$$\frac{\partial J}{\partial \theta_{i,k}} = x_i (y_k - g(x|\theta_{:,k})) = x_i (y_k - \hat{y}_k)$$

Multi-Class MLE + Logistic Activation

$$J = \sum_{k=1}^K y_k \ln(g(x|\theta_{:,k})) + (1 - y_k) \ln(1 - g(x|\theta_{:,k}))$$

$$\frac{\partial J}{\partial \theta_{i,k}} = x_i(y_k - \hat{y}_k)$$

- Vectorizing this for all i, k we get the matrix:

$$\frac{\partial J}{\partial \theta} = x^T(y - \hat{y})$$

- And the batch version would be

$$\frac{\partial J}{\partial \theta} = \frac{1}{N} X^T(Y - \hat{Y})$$

Multi-Class Evaluation

- Just like binary classification, we can evaluate the accuracy of a multi-class classifier:

$$accuracy = \frac{1}{N} \sum_{i=1}^N (Y_i == \hat{Y}_i)$$

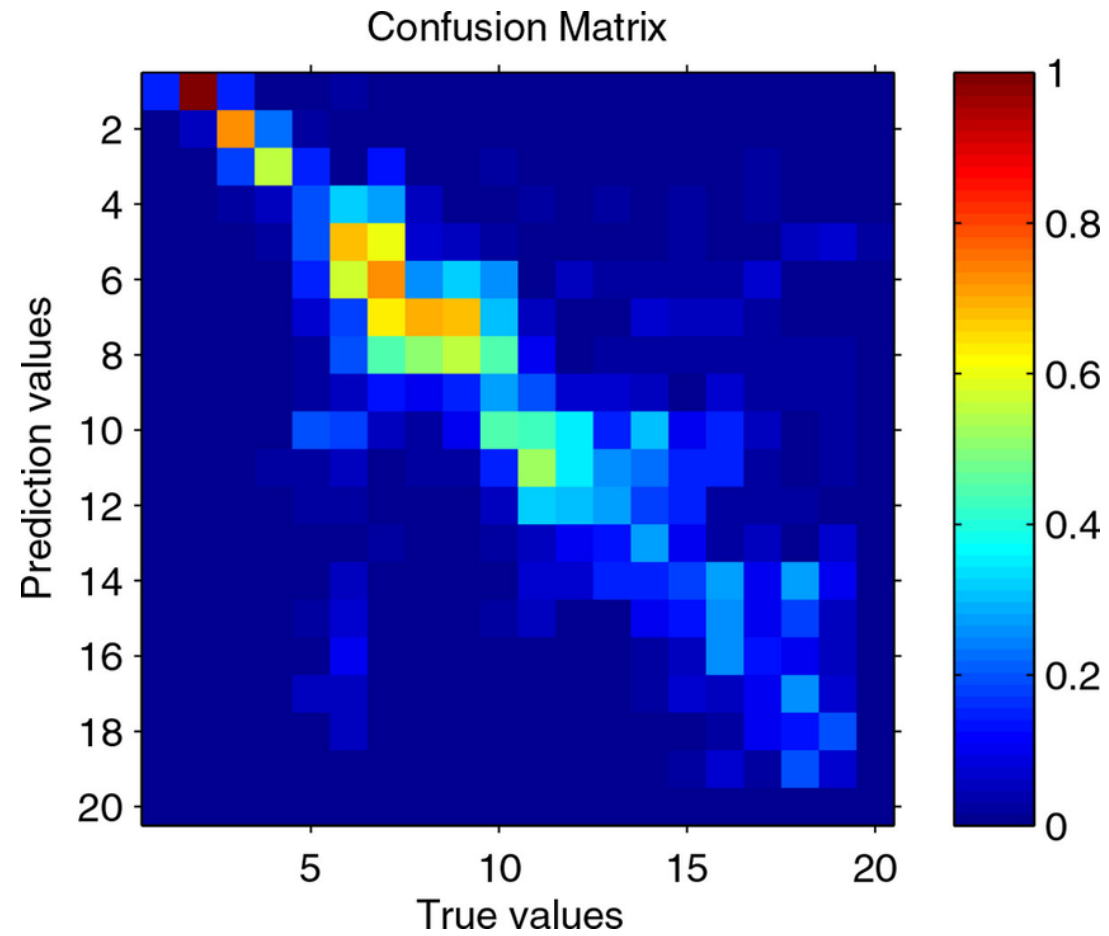
- In addition, particular to multi-class classification, we may be interested in investigating which classes get confused with which other classes
- To observe this, we can look at a *confusion matrix*

Confusion Matrix

All Confusion Matrix

Output Class	1	2	3	4	
	6670 20.8%	86 0.3%	1234 3.9%	558 1.7%	78.0% 22.0%
	28 0.1%	6476 20.2%	625 2.0%	804 2.5%	81.6% 18.4%
	801 2.5%	762 2.4%	5934 18.5%	127 0.4%	77.8% 22.2%
	501 1.6%	676 2.1%	207 0.6%	6511 20.3%	82.5% 17.5%
Target Class					
					1 2 3 4
					83.4% 16.6%
					81.0% 19.1%
					74.2% 25.8%
					81.4% 18.6%
					80.0% 20.0%

Confusion Matrix



Additional Objective Functions

- Another common objective function for multi-class classification is *cross entropy*
- So let's take a look at that.
- But to do this, we must first learn about the *softmax* activation function.

Softmax

- It is sometimes advantageous to be able to think of the output values as class *probabilities*.
- To be a valid distribution, we want these outputs to sum to one.
- The softmax function does this!
- The softmax function gives us back the probability of an observation belonging to class j , which we'll indicate as \hat{y}_j .
- This is computed as:

$$P(y = j) = \hat{y}_j = \frac{e^{x\theta_{:,j}}}{\sum_{k=1}^K e^{x\theta_{:,k}}}$$

Entropy

- Entropy measure the randomness in a system.
- Given a probability distribution $p = (p_1, p_2, \dots, p_K)$ such that $\sum_{k=1}^K p_k = 1$, the entropy is computed as:

$$H = \sum_{k=1}^K -p_k \ln p_k$$

Cross Entropy Loss

- If we have two distributions, a, b and want to compare them, we can use *cross entropy*:

$$H = \sum_{k=1}^K -a_k \ln b_k$$

- If we consider our target outputs to be y and our output is from the softmax activation function (or really any activation function that provides a valid output probability), then we can compute the ***cross-entropy loss objective function*** :

$$J = - \sum_{k=1}^K y_k \ln(\hat{y}_k)$$

Cross Entropy Loss

$$J = - \sum_{k=1}^K y_k \ln(\hat{y}_k)$$

- Since typically our y distribution will be all zeros, with just one location having a probability of one (say $y_a = 1$) then this simplifies to:

$$J = - \ln(\hat{y}_a)$$

Gradient of Softmax with Cross Entropy

- Now we need to find the update rule for all the parameters $\theta_{i,j}$, i.e. $\frac{\partial J}{\partial \theta_{i,j}}$, when we have a softmax function and cross entropy objective function.
- In order to do this, we need to write our cross-entropy objective function in terms of θ

$$J = -\ln(\hat{y}_a)$$
$$J = -\ln\left(\frac{e^{x\theta_{:,a}}}{\sum_{k=1}^K e^{x\theta_{:,k}}}\right)$$

Gradient of Cross Entropy

$$J = -\ln\left(\frac{e^{x\theta_{:,a}}}{\sum_{k=1}^K e^{x\theta_{:,k}}}\right)$$

- Now we can take the partial gradient of this with respect to a parameters $\theta_{i,j}$:

$$\frac{\partial J}{\partial \theta_{i,j}} = -\frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{\partial}{\partial \theta_{i,j}} \left(\frac{e^{x\theta_{:,a}}}{\sum_{k=1}^K e^{x\theta_{:,k}}} \right) \right)$$

- Now again there will be two cases

Gradient of Cross Entropy

$$\frac{\partial J}{\partial \theta_{i,j}} = - \frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{\partial}{\partial \theta_{i,j}} \left(\frac{e^{x\theta_{:,a}}}{\sum_{k=1}^K e^{x\theta_{:,k}}} \right) \right)$$

- Case 1: $j \neq a$

$$\frac{\partial J}{\partial \theta_{i,j}} = - \frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{(0 \sum_{k=1}^K e^{x\theta_{:,k}}) - (e^{x\theta_{:,a}} x_i e^{x\theta_{:,j}})}{(\sum_{k=1}^K e^{x\theta_{:,k}})^2} \right)$$

$$= \frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{e^{x\theta_{:,a}} x_i e^{x\theta_{:,j}}}{(\sum_{k=1}^K e^{x\theta_{:,k}})^2} \right) = \frac{x_i e^{x\theta_{:,j}}}{\sum_{k=1}^K e^{x\theta_{:,k}}} = x_i \hat{y}_j$$

Gradient of Cross Entropy

$$\frac{\partial J}{\partial \theta_{i,j}} = - \frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{\partial}{\partial \theta_{i,j}} \left(\frac{e^{x\theta_{:,a}}}{\sum_{k=1}^K e^{x\theta_{:,k}}} \right) \right)$$

- Case 2: $j = a$

$$\begin{aligned} \frac{\partial J}{\partial \theta_{i,j}} &= - \frac{\sum_{k=1}^K e^{x\theta_{:,k}}}{e^{x\theta_{:,a}}} \left(\frac{(x_i e^{x\theta_{:,j}} \sum_{k=1}^K e^{x\theta_{:,k}}) - (e^{x\theta_{:,j}} x_i e^{x\theta_{:,j}})}{(\sum_{k=1}^K e^{x\theta_{:,k}})^2} \right) \\ &= \frac{x_i (e^{x\theta_{:,j}} - \sum_{k=1}^K e^{x\theta_{:,k}})}{\sum_{k=1}^K e^{x\theta_{:,k}}} = x_i (\hat{y}_j - 1) \end{aligned}$$

- Thus:

$$\frac{\partial J}{\partial \theta_{i,j}} = \begin{cases} x_i \hat{y}_j & \text{if } j \neq a \\ x_i (\hat{y}_j - 1) & \text{Otherwise} \end{cases}$$

a: The right class.

Gradient of Cross Entropy

$$\frac{\partial J}{\partial \theta_{i,j}} = \begin{cases} x_i \hat{y}_j & \text{if } j \neq a \\ x_i (\hat{y}_j - 1) & \text{Otherwise} \end{cases}$$

- Since our target y is a binary vector, then we can write this as:

$$\frac{\partial J}{\partial \theta_{i,j}} = x_i (\hat{y}_j - y_j)$$

- If \hat{y} is the vector of predicted class probabilities, then we can vectorize this to update all the parameters at once as:

$$\frac{\partial J}{\partial \theta} = x^T (\hat{y} - y)$$

- And to do batch updates:

$$\frac{\partial J}{\partial \theta} = X^T (\hat{Y} - Y)$$

Cost_function



Linear regression

$$\frac{\partial J}{\partial \theta} = 2x^T(x\theta - y)$$

Multi_class with binary classifier

- Given all of this, let's adapt our **maximum** likelihood with logistic activation equations and gradient rules for multi-class classification:

$$\ell_{\theta}(x, y) = \prod_{k=1}^K g(x|\theta_{:,k})^{y_k} (1 - g(x|\theta_{:,k}))^{(1-y_k)}$$

$$\frac{\partial J}{\partial \theta} = \frac{1}{N} X^T (Y - \hat{Y})$$

Softmax with Cross Entropy

$$H = \sum_{k=1}^K -a_k \ln b_k$$

$$J = - \sum_{k=1}^K y_k \ln(\hat{y}_k)$$

$$\frac{\partial J}{\partial \theta_{i,j}} = x_i(\hat{y}_j - y_j)$$

- Since typically our y distribution will be all zeros, with just one location having a probability of one (say $y_a = 1$) then this simplifies to:

$$J = -\ln(\hat{y}_a)$$

$$\frac{\partial J}{\partial \theta_{i,j}} = \begin{cases} x_i \hat{y}_j & \text{if } j \neq a \\ x_i(\hat{y}_j - 1) & \text{Otherwise} \end{cases}$$

a: The right class.