

CS 615 – Deep Learning

Generative (Adversarial) Networks

Generative Models


- With RNNs and LSTMs we talked about how a one-to-many system can be used as a *generative model*
 - Given an initial seed/input, x , generate outputs $y^{(1)}, y^{(2)}, \dots, y^{(T)}$
 - For example, generating text based on a seed input.
- Another approach to creating a model that can generate “realistic” output, is to create a *generative adversarial network (GAN)*.

Generative Adversarial Networks (GAN)

- A Generative Adversarial Network (GAN) uses a game-theoretic approach to learn from a training distribution
- It is a framework proposed by Ian Goodfellow, Yoshua Bengio and others in 2014
- It is a hot topic of discussion and research these days because of its high potential

GAN

- Basic idea of a GAN:
 - It is a “game” between two players: a generator (G) and a discriminator (D)
 - The generator tries to fool the discriminator by generating realistic-looking data
 - The discriminator tries to distinguish between real (training set) and fake (generated) data



GAN

- The goal of the generator is to learn how to generate data that has a similar distribution to the real data x
- We will call $G(z)$ our *generator network*
 - It takes as its input a random feature vector, z , generated based on the training data's input distribution.
 - Its output will be some generated data $G(z)$
- Similarly, let $D(x)$ be our *discriminator network*
 - This takes as an input either a real observation, x , or one generated by the generator network, i.e. $G(z)$.
 - This network returns the probability that it came from a real observation as opposed to having been generated by $G(z)$

GAN Objective Functions

- We train D and G simultaneously by playing a **two-player minimax game**
- Let's imagine that we feed our discriminative network several sets of real observations and fake/generated ones.
- Doing this over several “observations”, since $D(x)$ returns a probability, we can define a likelihood objective function for our discriminant function:

$$J_d = \prod_{i=1}^N D(X_i) (1 - D(G(Z_i)))$$

- Or we can take the log of this to get a log likelihood objective function:

$$J_d = \sum_{i=1}^N (\log(D(X_i)) + \log(1 - D(G(Z_i))))$$

GAN Objective Functions

$$J_d = \sum_{i=1}^N (\log(D(X_i)) + \log(1 - D(G(Z_i))))$$

- So the discriminator wants to *maximize this*.
- The generative network wants to trick the discriminator.
- So its objective function for the generative network could be:

$$J_g = \sum_{i=1}^N (\log(D(G(Z_i))))$$

Training GAN

- Just alternate between optimizing J_d and optimizing J_g using all the data (full batch)
 - Can be computationally prohibitive
 - May result in overfitting for finite datasets.
- So instead, in practice, we alternate between k steps of optimizing J_d with mini-batches and one step of optimizing J_g

Pseudocode

1. Obtain the distribution p_z based on training data.
2. Until termination criteria is met
 1. For k steps
 1. Create m (half of mini-batch) “fake” input vectors based on the distribution p_z .
 2. Grab m (other half of mini-batch) true input vectors from the training data.
 3. Update/train D using this mini-batch
 2. Generate m (half of mini-batch) “fake” input vectors based on distribution p_z .
 3. Grab m (other half of mini-batch) true input vectors from the training data.
 4. Update/train G using this mini-batch

Generated Samples




Fig1: MNIST dataset (Ian Goodfellow)




Fig2: TFD dataset (Ian Goodfellow)

The rightmost column shows the nearest training example of the neighboring sample

Bibliography

- Goodfellow, Ian, et al. "Generative adversarial nets." *Advances in neural information processing systems*. 2014.
- <https://towardsdatascience.com/understanding-generative-adversarial-networks-gans-cd6e4651a29>